
The Autonomous Enterprise

A Multi-Agent AI System for Sharp Reasoning and
Independent Decision-Making

The Vision: Next-Generation Workforce

- **Sales Engineer Agent**

Technical qualification, solution configuration, and pre-sales support to accelerate sales cycles.

- **Customer Receptionist Agent**

First-line interaction, intelligent routing, and personalized welcome to enhance customer experience.

- **Account Executive Agent**

Deal finalization, negotiation, pricing optimization, and contract closure to drive revenue growth.

- **Tech Support Agent**

Tier 1 issue resolution, knowledge base querying, and intelligent ticket escalation.



Core Design Principles

Engineering foundations for technical strength and real-world impact through four core principles, powered by Google GenAI and Google ADK agents framework.

Sharp Reasoning

Google GenAI LLM models deliver advanced reasoning chains to solve complex problems with precision.

Independent Decision-Making

Google ADK agents framework empowers agents to execute multi-step workflows without constant intervention.

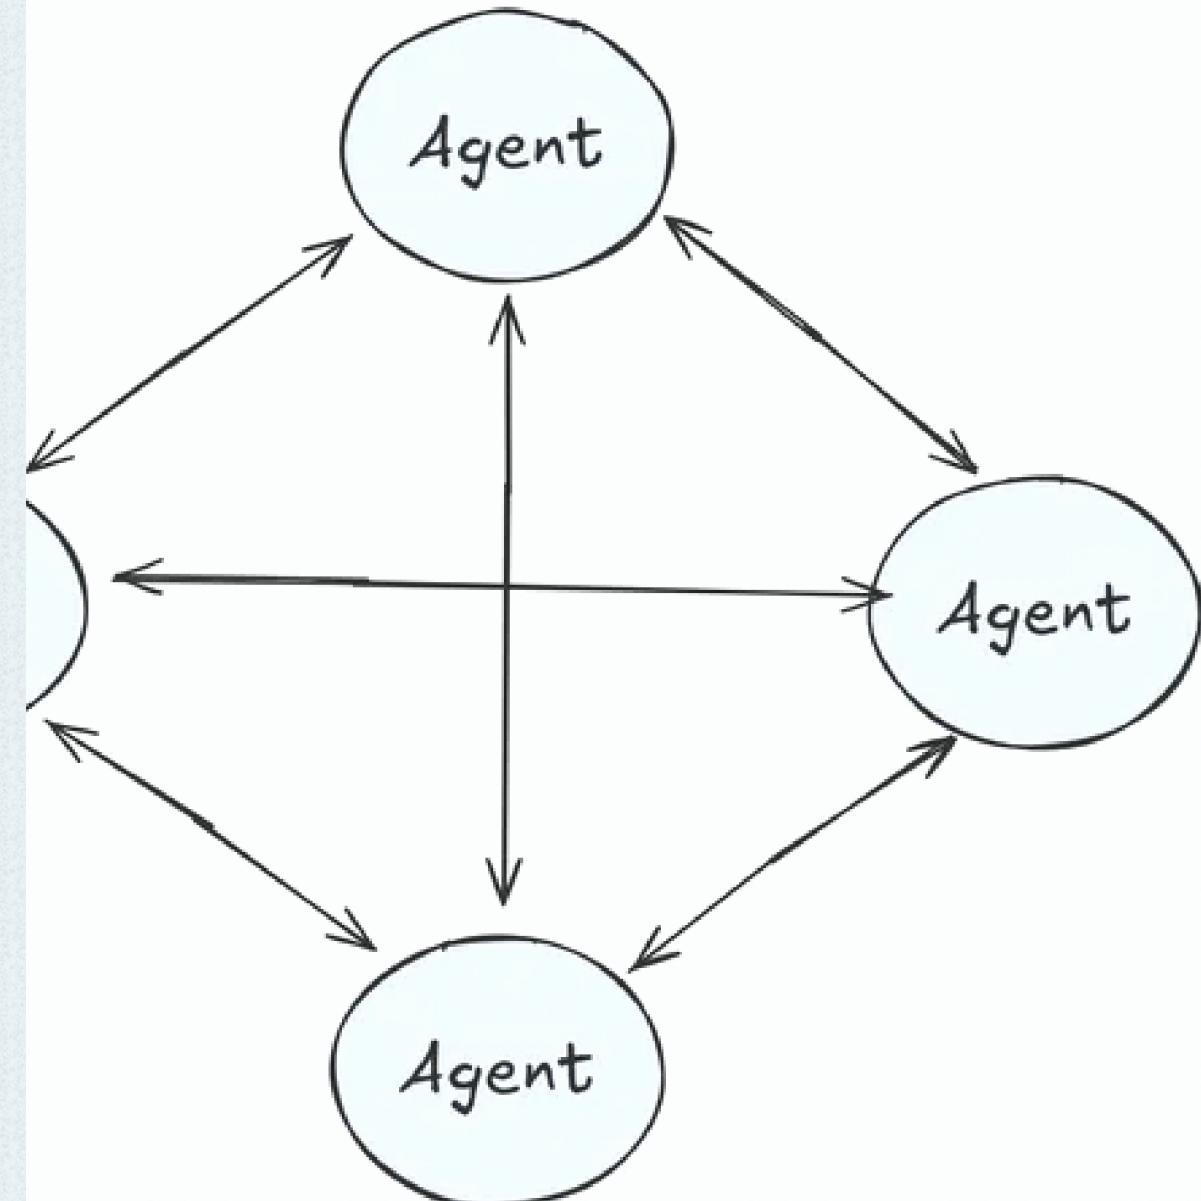
Safe Tool Integration

Seamless and secure integration with industry-relevant tools and external systems through MCP.

Originality & Impact

Deliver solutions demonstrating technical innovation and measurable real-world business impact.

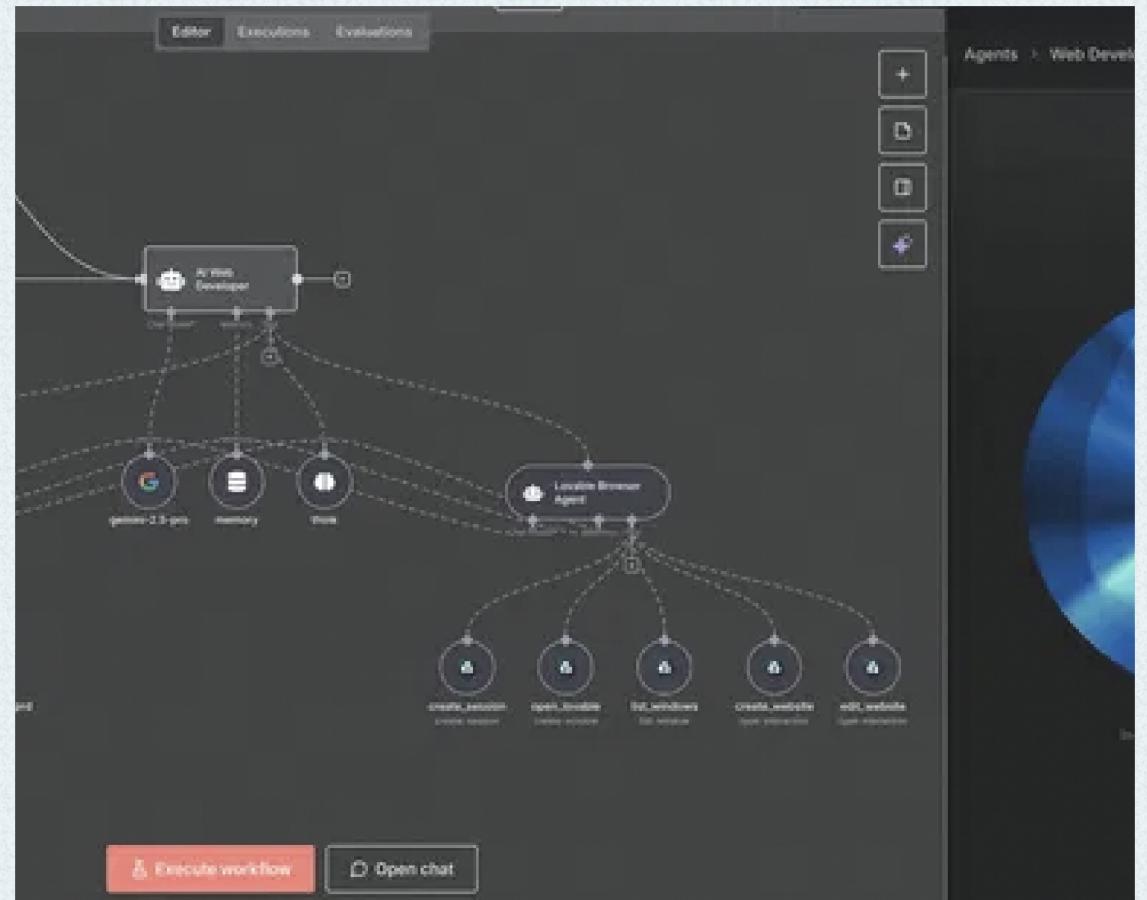
Network Architecture



Hyper-Realistic Voice & Interaction

ElevenLabs voice agents deliver highly expressive, low-latency, and natural conversational experiences that feel genuinely human, transforming customer interactions.

- 🔊 Human-like interaction with natural voice synthesis and expressive communication
- 👤 Enhanced customer reception through professional, engaging first impressions
- ⌚ Voice-enabled workflows enabling faster, more intuitive service delivery



Secure and Rapid Deployment

Daytona provides agent-native infrastructure with enterprise-grade security, enabling rapid deployment and reliable execution of AI agents.

Secure Execution

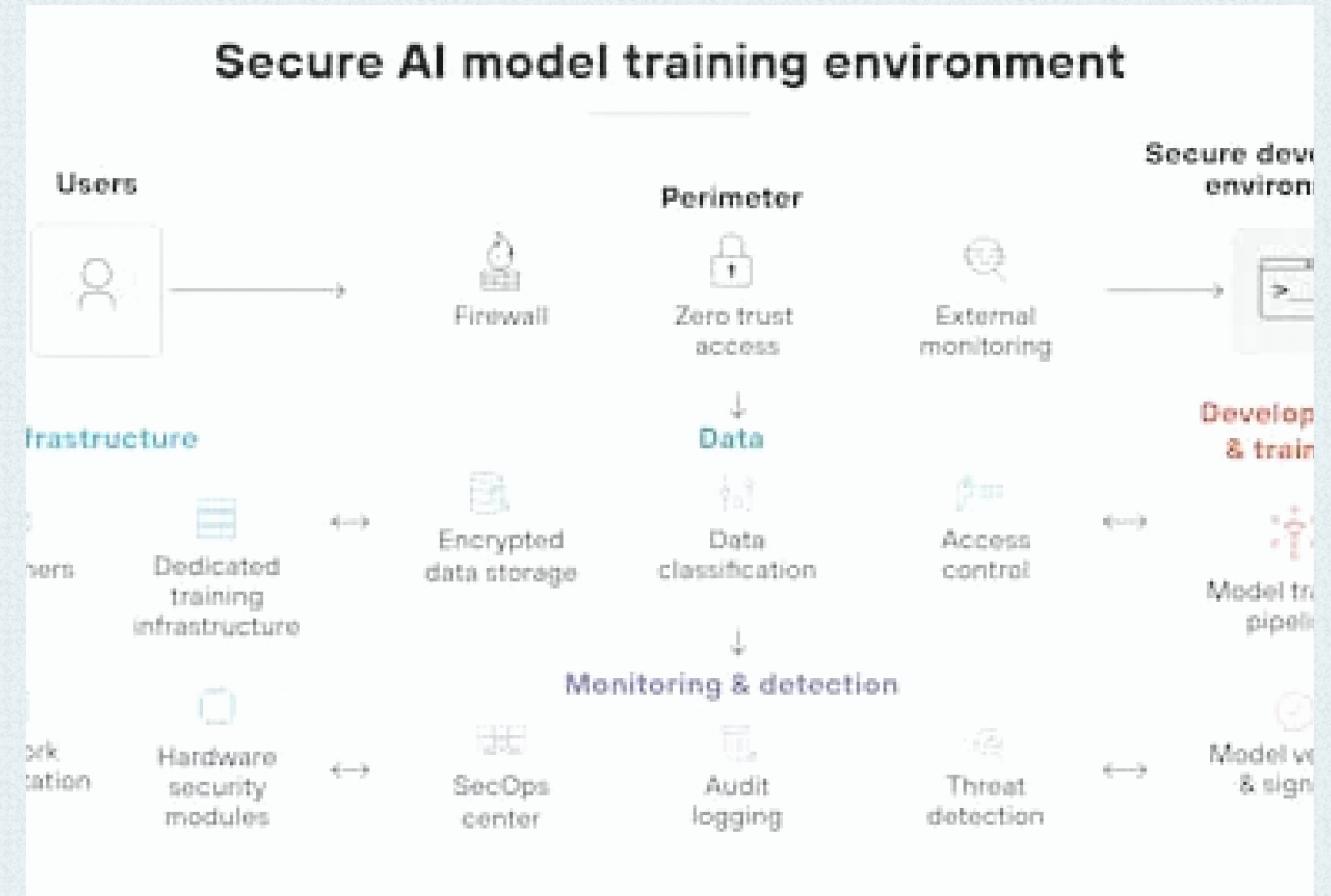
Isolated, ephemeral environments mitigate risks of running untrusted code

Rapid Creation

90ms environment creation supports agile development and rapid iteration

Stateful Operations

Persistence for agents to maintain context across complex multi-step tasks



Integrated Quality and Safety

A comprehensive safeguard framework combining Google GenAI, Google ADK, CodeRabbit, Galileo, and Daytona for intelligent, reliable, and observable agent operations.

Google GenAI & ADK Foundation

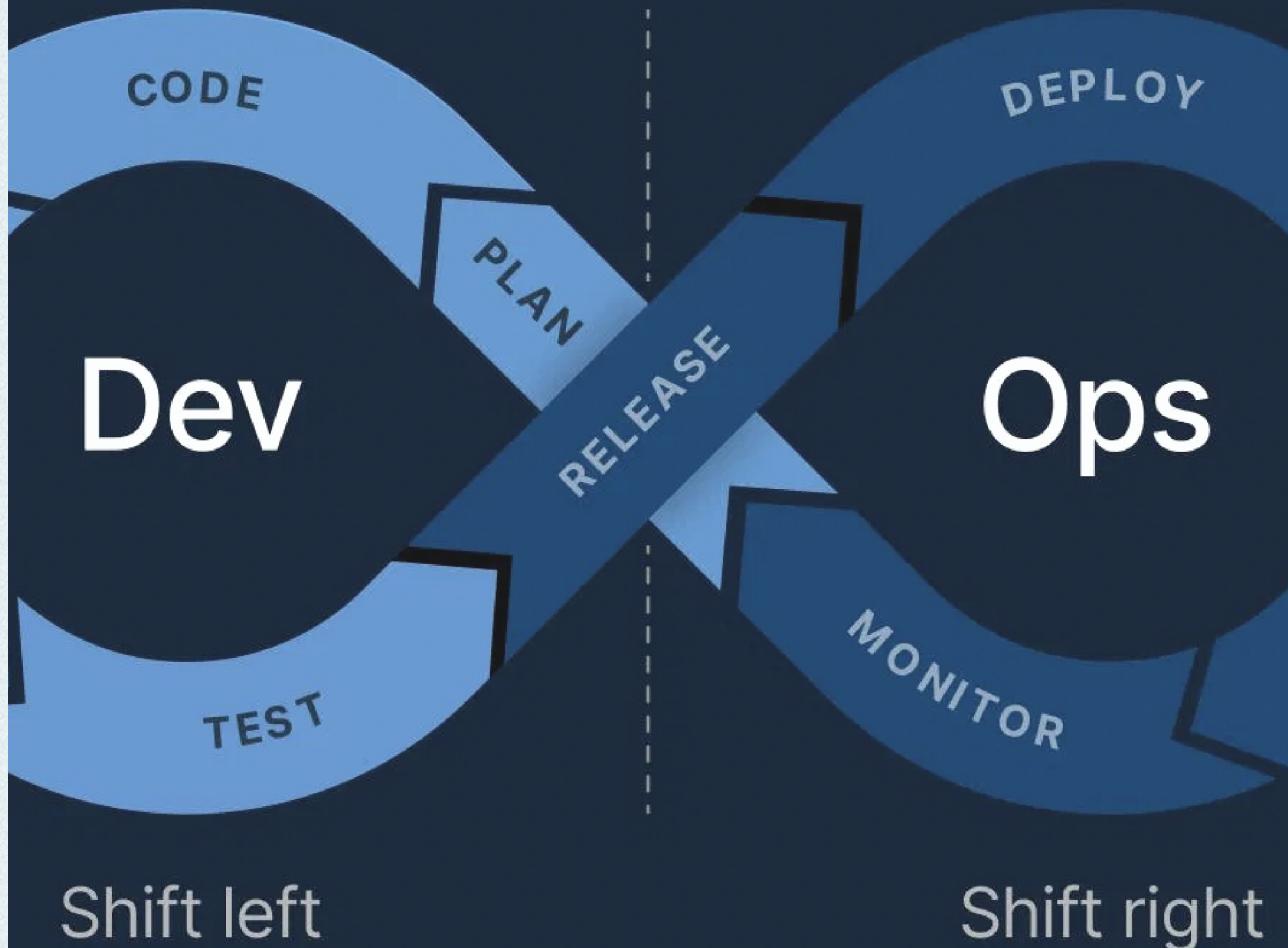
Sharp reasoning and independent decision-making powered by Google's advanced LLM models and agent framework

CodeRabbit MCP Code Review

Automated AI-powered code reviews ensuring agentic logic is robust, secure, and free of vulnerabilities

Galileo Observability & Daytona Deployment

End-to-end tracing and secure execution with continuous monitoring for operational excellence



Observability and Traceability

Galileo provides comprehensive observability and tracking for voice agents, enabling end-to-end visibility into agent behavior and performance.

End-to-End Tracing

Track voice agent interactions from customer input through decision-making to action execution

Debugging Non-Determinism

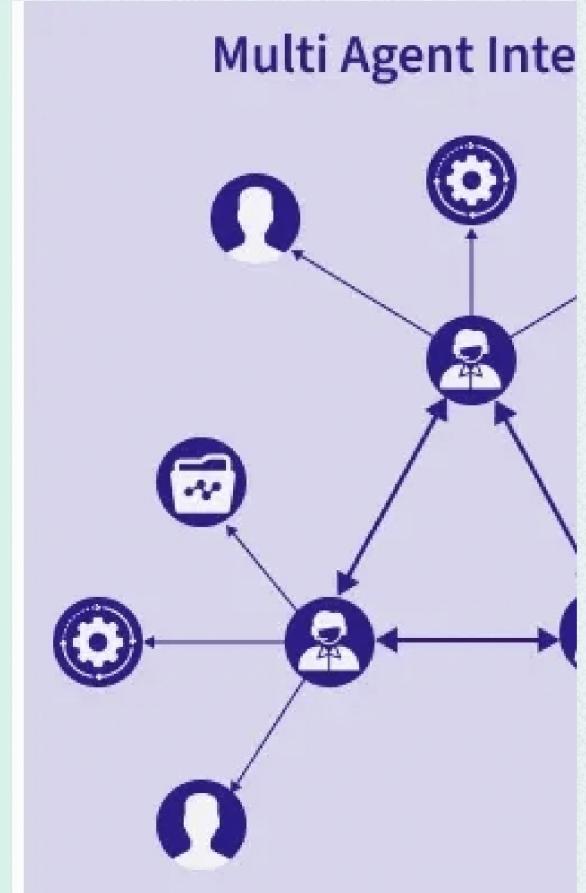
Debug the non-deterministic nature of LLM-powered voice agents to prevent failures

Guardrails and Evaluation

Set performance guardrails and evaluate agent behavior against expected outcomes

Operational Insight

Gain deep insight into voice agent performance, cost management, and reasoning quality



The Multi-Agent Architecture in Action

A seamless orchestration of specialized agents powered by Google GenAI LLM models, Google ADK agents framework, ElevenLabs voice, Daytona deployment, CodeRabbit MCP code review, and Galileo observability.

Customer Interaction

Receptionist Agent receives calls via ElevenLabs voice, identifies intent, and routes to specialists

Specialist Execution

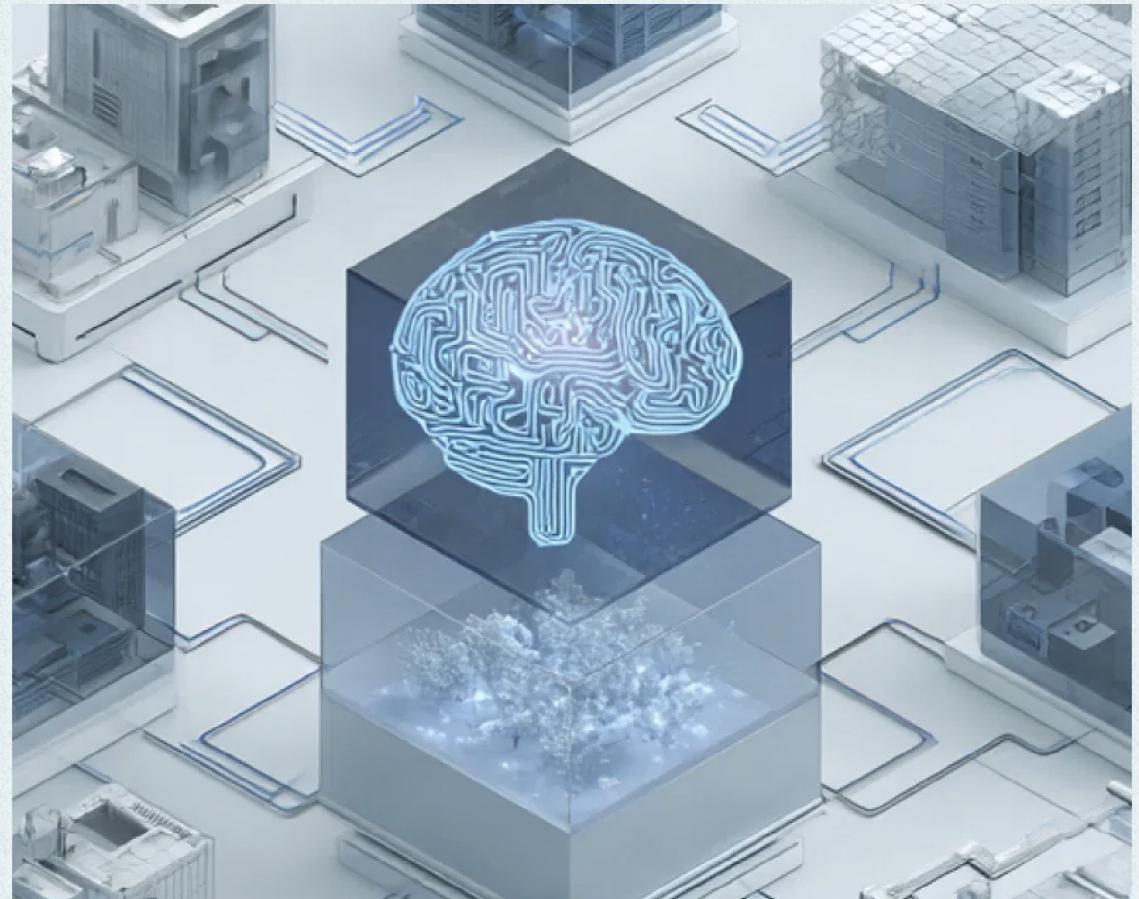
Sales Engineer qualifies opportunities, Account Executive negotiates and closes deals, Tech Support resolves issues securely on Daytona

Continuous Feedback Loop

CodeRabbit MCP ensures code quality, Galileo tracks performance and feeds insights

Seamless Outcome

Secure, highly effective automated service delivery powered by Google GenAI at enterprise scale



Measuring Success and Real-World Impact

Key performance indicators that measure the effectiveness and business value of the multi-agent system.

Efficiency (AHT)

Reduction in average handle time per customer interaction across all agents

Quality (FCR)

Increase in first contact resolution rate for Tech Support and customer satisfaction

Revenue (SQL)

Improvement in Sales Qualified Lead conversion rate by the Sales Engineer Agent

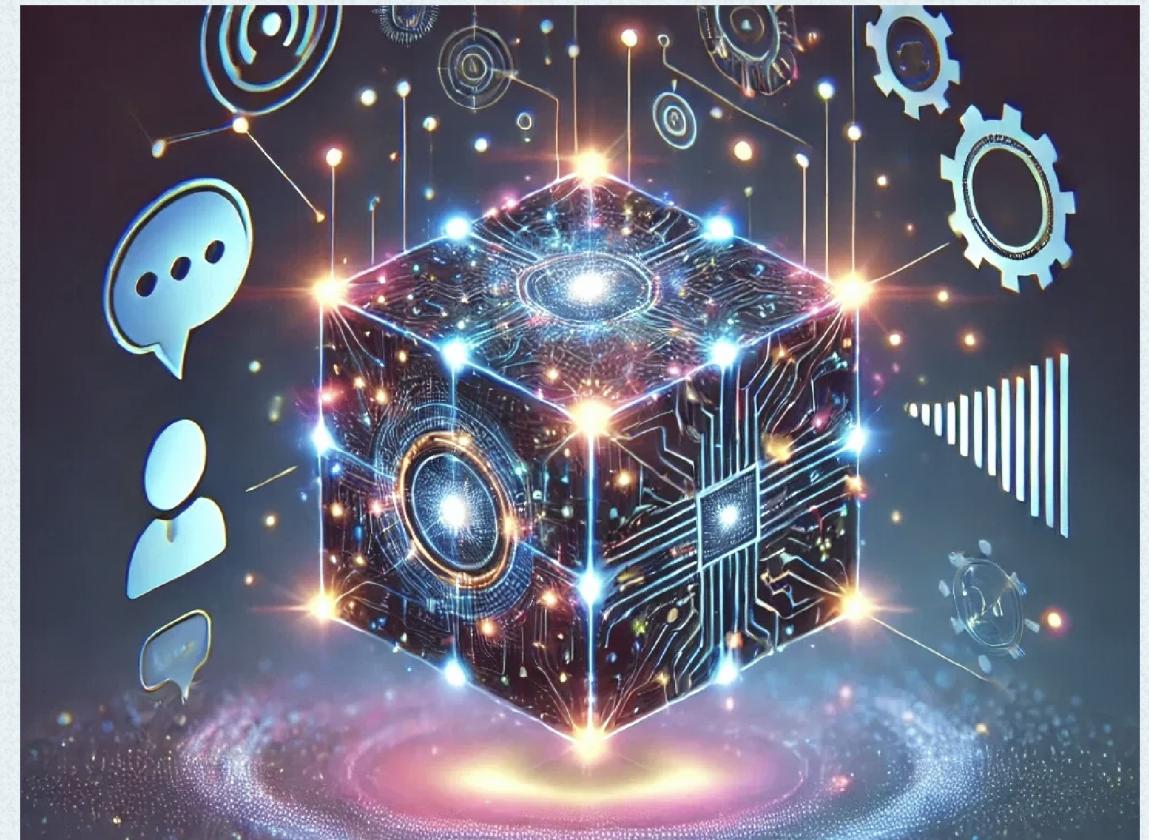
Stability (Reliability)

Reduction in production bugs and agent failures monitored by CodeRabbit and Galileo

Building the Autonomous Future, Responsibly

The multi-agent system powered by Google GenAI LLM models and Google ADK agents framework represents a transformative approach to enterprise automation—delivering originality, technical strength, and measurable real-world impact through sharp reasoning, independent decision-making, and safe tool integration.

- 💡 Embrace Google GenAI and ADK framework for intelligent, autonomous agent orchestration
- ⚙️ Master the integrated stack: ElevenLabs voice, Daytona, CodeRabbit MCP, and Galileo tracking
- 🛡️ Commit to continuous improvement through observability, code review, and safeguards



"The future of enterprise automation is intelligent, autonomous, and responsible."

The Journey: From Vision to Reality

Inspiration

Enterprise sales and support teams struggle with workload capacity, customer experience consistency, and response times. We envisioned intelligent agents handling the entire customer lifecycle: reception, technical qualification, deal negotiation, and support—freeing human teams to focus on complex, high-value relationships.

What We Learned

- Sharp reasoning requires careful prompt engineering and model selection
- Safeguards (code review, observability) are essential, not optional
- Voice agents need low-latency, expressive interactions for user trust
- Continuous monitoring and feedback loops drive continuous improvement

Challenges

- Building reliable, non-deterministic LLM-powered agents at scale
- Ensuring security and safe execution in production environments
- Achieving end-to-end observability and debugging agent behavior
- Integrating voice interactions with complex reasoning and negotiation workflows

What's Next

- Phase 1: Deploy Sales Engineer, Account Executive, and Receptionist agents
- Phase 2: Expand to Tech Support and specialized domain agents
- Phase 3: Integrate with CRM, ticketing, and knowledge base systems
- Continuous: Improve reasoning, expand capabilities, scale globally